

Close *the loop.*

Prove donation impact. Measure participant wellbeing. One platform, one source of truth.



Joshua Farrugia

Black Nova VC



Yucca Reinecke

EY · AI for Good



Anthony Hook

Product & marketing tech



Alejandra Castro

Data analyst



FOLLOW ALONG

Scan to open the deployed app. /presentation has the live in-product deck.

THE SHAPE OF THE WORK

Two stories Too Good can't yet tell.

80h

to assemble a single quarterly impact report. Pulled by hand from CSnet, Tanda and spreadsheets across sites.

Source: Rob Caslick, NFP briefing call, 29 May 2026

0

donors who hear what their meal actually did. Hundreds give. No system connects a gift to the shelter it served.

Source: TG-2 problem statement, G'dAI brief

BUILT BEFORE THE DAY

Not a vibe-coded mockup. A working app.

FRONTEND

Next.js 15 App Router · React Server Components · shadcn/ui

DATA

Drizzle ORM | Cloudflare D1 SQLite

AI

Cloudflare Workers AI · gpt-oss-120b · brand-voice grounded

DEPLOY

Cloudflare Workers via OpenNext · edge runtime · runs cheap

QUALITY

Vitest suite over import, scoring, validation, privacy, matching

GOVERNANCE

4 roles (admin, impact, ops, viewer) · audit log · frozen snapshots

Two surfaces live today

- /wellbeing
- /donations
- /wellbeing/reports
- /donations/ledger
- /import
- /donations/queue
- /ai
- /donations/shelters

Wellbeing imports the PWI Client Tracker, validates, scores and freezes a reproducible quarterly snapshot. Donations holds a unified ledger and a sensitive-shelter-safe review queue.


PWI tracker in. Reproducible report out.

Upload the tracker, validate the data quality, generate a frozen snapshot that won't drift if the source changes. Two grounded AI assistants share the same brand voice and never invent figures.

- 14 metric definitions (PWI + work-readiness extensions), configurable, scale-aware
- Cohort breakdown, change scores, validation issues, completeness per time point
- Two AI assistants: dashboard-level (across reports) and per-report, both grounded on the snapshot
- Six Recharts surfaces · CSV and XLSX export from the same frozen snapshot

Personal Wellbeing Index — domain movement

Average satisfaction (0–10) at baseline, 3 months and 6 months.



Donor. Order. Shelter. One traceable line.



MATCH METHOD

Postcode + product + quantity

BELOW THRESHOLD

Goes to manual review queue

SENSITIVE SHELTERS

Name & location masked

FROM SHIPPED TO ASPIRED

Built means the team can *generate* the report.
It doesn't yet mean they can
ask the right questions of it.

Here's where each PRD takes that further.

ALIGNED TO THE PROBLEM · *walk it live at /future/wellness*

From snapshot to story.

- 01 CSnet ingestion, alongside the PWI tracker
Parser tested against the real export. Long-format storage, idempotent, dedupes the same survey across multiple comparison tabs.
- 02 Longitudinal participant master
Stable CSnet client ID. One woman, tracked across cohorts and years. Cohort 24 vs 26, side by side, in seconds.
- 03 Tanda for attendance and wages
Mocked for the demo, scoped for the API connector after. Hours by site, lateness, wages paid. Replaces the monthly hand-pull.
- 04 Macro context overlay
Program changes and external indicators on the same timeline as outcome trends. The 'why' alongside the 'what'.

FROM THE LEAD

"She doesn't trust the system or the reports CSnet delivers. It's just raw data she can't make sense of."

Notes · Two Good briefing · 29 May 2026

IN THE APP, RIGHT NOW

Vision-mode page at */future/wellness* walks each of these.

HOUSE RULE, BAKED IN

AI-drafted narrative is always a draft. The lead reviews. Never auto-sent.


From "we suspect" to "here's the evidence".

Pick a metric. Pick an attribute. The view returns subgroup performance with sample-size labels and $n < 5$ suppression. Ethical by construction.

- Cohort, site, age band, housing at intake, country of origin, language, faith. Consented at intake, never required.
- Never exposed in donor-facing or public reporting.
- "Why?" returns a brand-voiced hypothesis grounded in the context overlay. Not a causal claim.

PWI total · 6mo gain · by housing at intake

Synthetic · Intakes
23-26



Participants with insecure housing at intake show a 2.5-point lower 6-month PWI gain than the program average. The context overlay flags three rental affordability events in their cohort's window.

Three sources, one ledger, one trace.

01

Unified ledger, two pools

Meals and care packs never mix. Each has its own donor pool, ledger, carry-forward balance. Live at /donations/ledger and /donations-beta.

02

Weekly allocation engine

Like-for-like demand by product subtype. Priority: exact quantity match → largest donor orders first → Too Good gap. Carry-forward balances roll over.

03

Manual review queue

Below-threshold matches surface to ops. Sensitive shelters masked from anyone without the permission.

04

Donor impact reports · shipped

Aggregated per email. Printable PDF report pages. Delivered via Resend with a link. Live at /donations-beta.

PRIVACY BY DESIGN

Order IDs across customer and shelter systems are never directly compared. Sensitive shelters masked in every partner-facing export. Each trace stores its method, confidence and the candidates considered.

COVERAGE TODAY

Vitest over match scoring, sensitive-shelter privacy, donor normalisation, meal donation matching.

LIVE DEMO · /DONATIONS-BETA · WEEKLY ALLOCATION ENGINE

Donor to shelter, every week, on the record.

- 01** Weekly allocation engine
Meals and care packs run separately. 11 weekly cycles matched across 2026, carry-forward balances preserved each week.
- 02** Gap tracking & carry-forward
1,110 meal gap flagged for marketing. 648 care pack units rolled forward. No allocation is lost — every shortfall is visible and named.
- 03** Donor impact reports · shipped
Per-donor allocation ledger, 354 rows. Printable PDF reports. Delivered via Resend at /donations-beta/reports.

MEALS ALLOCATED

1,536

TOO GOOD GAP

1,110

[→ start marketing](#)

LEDGER ROWS

354 rows · 11 weeks · imported from Close the Loop deidentified data

ALIGNED TO THE PROBLEM • WALK IT LIVE AT /FUTURE/IMPACT

Generic thanks ➤ real connection.

Hundreds of supporters give. The meals reach women in crisis. The donor hears nothing meaningful again. The relationship cools, the next ask works harder.

TODAY

"Thanks for your generous donation."

PRD2

"Your gift reached Shelter X on Tuesday. It funded 24 meals for women starting over."

A draft, in their voice, for a human to send.

The AI assistant drafts the thank-you in Two Good's brand voice, grounded only on the matched trace. The team reviews, edits, sends. No auto-send.

- Brand voice profile injected: warm, human, purposeful, dignity-led
- Citation-only: every figure traceable to the underlying donation match
- Resend wired for delivery. Existing email template, brand-coloured
- Foundation for re-engagement, not a one-off transaction

FROM Two Good Co <hello@twogood.com.au>
TO donor@example.com
SUBJECT Where your meal reached, with gratitude.

Hi Sam,

When you bought **two Buy One Give One meals** on 12 April, you funded two meals that reached a partner shelter in **Western Sydney** on 18 April.

Women starting over were warm, and fed, that night, in part because of you.

— Drafted by the assistant, grounded in trace ID TRC-0412-WS. Reviewed by Imogen before send.

REVIEW & SEND

REGENERATE

SAVE AS TEMPLATE

QUALITY OF PRESENTATION · ALIGNED TO ALL THREE JUDGING CRITERIA

Built for the team of one.

FOR THE IMPACT LENS

"Will it have *real impact*?"

80 hours/quarter → minutes

Equity questions become askable

Donors stop falling off the back of the truck

FOR THE TECHNICAL LENS

"Will it actually *work*?"

Deployed, not mocked. Edge runtime.

Test suite over ingestion, scoring, privacy, matching
No keys client-side. Audit log on every action.

FOR THE ADOPTION LENS

"Will the NFP *actually adopt it*?"

No data analyst on staff required

No ML PhD, no migrations to babysit

Cloudflare cost ≈ a few coffees per month at this scale

IN CLOSING

Close the *loop.*

For the women in the program, the donors who fund it,
and the team of one who has to make sense of it all.

BUY GOOD • DO GOOD



TAKE THE APP WITH YOU

Scan to explore /presentation,
/future/wellness, /future/impact in the
deployed app.

80 hours » minutes

The quarterly report, generated and frozen.

"We suspect" » "Here's the evidence"

Equity questions, with sample-size honesty.

Generic thanks » real connection

Every donor, tied to the shelter they reached.